

Measuring Regression Performance

- Mean Absolute Error (MAE)

```
#Firstly, we normalise our dataset to a common scale using the min max scaler
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
normalised_df = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)
features_df = normalised_df.drop(columns=['Heating_Load', 'Cooling_Load'])
heating_target = normalised_df['Heating_Load']
```

#Now, we split our dataset into the training and testing dataset. Recall that we had earlier segmented the features and target variables.

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(features_df, heating_target,
test_size=0.3, random_state=1)
```

```
linear_model = LinearRegression()
#fit the model to the training dataset
linear_model.fit(x_train, y_train)
#obtain predictions
predicted_values = linear_model.predict(x_test)
#MAE
from sklearn.metrics import mean_absolute_error
mae = mean_absolute_error(y_test, predicted_values)
round(mae, 3)      #prints 0.063
```

- Residual Sum of Squares (RSS)

```
import numpy as np
rss = np.sum(np.square(y_test - predicted_values))
round(rss, 3)      #prints 1.823
```

- Root Mean Square Error (RMSE)

```
from sklearn.metrics import mean_squared_error
rmse = np.sqrt(mean_squared_error(y_test, predicted_values))
round(rmse, 3)      #prints 0.089
```

- R-Squared

```
from sklearn.metrics import r2_score
r2_score = r2_score(y_test, predicted_values)
round(r2_score, 3)  #prints 0.893
```